

MATHEMATICS OF REINFORCEMENT LEARNING WITH APPLICATIONS TO QUANTITATIVE FINANCE

Yuanhao JIANG^{1,†,*}, Luis Hernández^{2,†}, Catherine Parr^{3,*}

¹The University of Edinburgh, Edinburgh, UK

²Institution or Company Name, City, State

³Hampton Court Palace, Richmond, England

ABSTRACT

This paper is an example of and a \LaTeX template for typesetting ASME conference papers using the `asmeconf` class. This \LaTeX template follows ASME guidelines for margins, fonts, headings, captions, and reference formats as of 2022. The class is intended to be used with the `asmeconf.bst` \BibTeX style for reference formatting, which is part of this distribution. The template produces pdfs that contain hyperlinks, bookmarks, and metadata; and references can include the DOI and URL fields. Links may be colored, for online use, or black, for publication. The class enables inline author names, following ASME's current style, but can also produce the traditional grid style. Options include line numbering, final column balancing, various math options, government copyright, and archivability (PDF/A). In addition, section headers may contain mathematics, references, citations, and footnotes. The class is compatible with $\pdf\LaTeX$ or $\Lua\LaTeX$.

Keywords: Reinforcement Learning, Artificial Neural Networks

1. INTRODUCTION

The object is to use reinforcement learning to solve the following scenario: a customer arrives with a number of features, we want to promote the customer with a proper price, such that he is most likely to purchase our product, and we also have the highest profit and a better portfolio.

We can consider this scenario as markov decision process, with the current state being the current customer features and current portfolio. At each time step, the agent promote a price according to the current customer features and the current portfolio, after that the customer resopnse to the action. The reward is then given based on the customer response and the portfolio. Then the portfolio updates and a new customer arrives.




FIGURE 1: MDP ILLUSTRATION

2. REINFORCEMENT LEARNING

The reinforcement learning method we want to use is based on the Markov decision process (MDP), as illustrated in Figure 1. However, in RL, the agent is not told which actions to take, but instead must discover which actions yield the most reward by trying them. Our goal at each time step is to find a policy π to maximize the expected reward along the process afterwards, which is given by the value function

$$V(s) = \mathbb{E} \left[\sum_{i=0}^{\infty} \gamma^i R_{t+i+1} \mid S_t = s \right]$$

To achieve this, we need to evaluate the current policy, and improve it accordingly.

2.1 Policy Evaluation

We want to evaluate the value function under a given policy π . According to Bellman equation,

$$v_{\pi}(s) = \mathbb{E}_{\pi} [R_{t+1} + \gamma v_{\pi}(S_{t+1}) \mid S_t = s]$$

We can use it to iteratively compute v_{π} . The iteratively update is

$$v_{k+1}(s) = \mathbb{E}_{\pi} [R_{t+1} + \gamma v_k(S_{t+1}) \mid S_t = s]$$

[†]Joint first authors

*Corresponding author: s2132254@ed.ac.uk, kate@thepalace.gov
Documentation for `asmeconf.cls`: Version 1.31, August 18, 2022.

which is proven that it will converge to the true value given any initial value function.

2.2 Policy Improvement

Now we have the evaluated value function for the policy, we want to improve it accordingly. To do this, we introduce the q-function under policy π

$$q_\pi(s, a) = \mathbb{E} [R_{t+1} + \gamma v_\pi(S_{t+1}) \mid S_t = s, A_t = a]$$

The policy improvement theorem says that, let π and π' be any pair of policy, for all s we have

$$q_\pi(s, \pi'(s)) \geq v_\pi(s) \Rightarrow v_{\pi'}(s) \geq v_\pi(s)$$

Now we can improve the policy by

$$\pi'(s) = \underset{a}{\operatorname{argmax}} q_\pi(s, a)$$

for all s .

2.3 Policy Iteration

By iteratively applying policy evaluation and policy improvement, we can obtain a sequence of monotonically improving policies and value functions:

$$\pi_0 \xrightarrow{E} v_{\pi_0} \xrightarrow{I} \pi_1 \xrightarrow{E} v_{\pi_1} \xrightarrow{I} \dots \xrightarrow{I} \pi_* \xrightarrow{E} v_*$$

where \xrightarrow{E} denotes a policy evaluation and \xrightarrow{I} denotes a policy improvement. This way of finding an optimal policy is called policy iteration.

2.4 Value Iteration

One drawback to policy iteration is that each of its iterations involves policy evaluation, which itself is a protracted iterative computation requiring multiple sweeps through the state set.

We now truncate the policy evaluation such that it is stopped after just one sweep (one update of each state):

$$\begin{aligned} v_{k+1}(s) &= \max_a q(s, a) \\ &= \max_a \mathbb{E} [R_{t+1} + \gamma v_k(S_{t+1}) \mid S_t = s, A_t = a] \end{aligned}$$

We keep iterate until v converges (to v_*), then we record the last policy $\pi(s) = \underset{a}{\operatorname{argmax}} q(s, a)$.

2.5 Generalized Policy Iteration

As long as both processes, policy evaluation and policy improvement, continue to update all states, the ultimate result is typically the same — converge to the optimal value function and an optimal policy.

Generalized policy iteration (GPI) is used to refer to the general idea of letting policy evaluation and policy improvement interact, independent of the granularity of the two processes (instead of letting each complete before the other begins).

2.6 Temporal-Difference Learning

When we don't have the complete model of the environment, we need to estimate the value function from only experience. There are two methods to do so: Monte Carlo Methods and Temporal-Difference. We will use Temporal-Difference method in this project.

To estimate $v_\pi(S_t)$ using only the experience, we take the average on all experienced trajectory' values. Temporal-Difference method uses $G_i = [R_{t+1}]_i + \gamma[v_\pi(S_{t+1})]_i$ to estimate the value of the i th trajectory starting at S_t , with v_π being the estimated value function at the i th time S_t is visited, which can be initialized to anything at start. Now we have the incremental formula to update our value function:

$$\begin{aligned} [v_\pi(S_t)]_{k+1} &= \frac{1}{k} \sum_{i=1}^k G_i \\ &= \frac{1}{k} \left(G_k + \sum_{i=1}^{k-1} G_i \right) \\ &= \frac{1}{k} (G_k + (k-1)[v_\pi(S_t)]_k) \\ &= [v_\pi(S_t)]_k + \frac{1}{k} (G_k - [v_\pi(S_t)]_k) \\ &= [v_\pi(S_t)]_k + \frac{1}{k} ([R_{t+1}]_k + \gamma[v_\pi(S_{t+1})]_k - [v_\pi(S_t)]_k) \end{aligned}$$

To simplify:

$$v_\pi(S_t) \leftarrow v_\pi(S_t) + \alpha [R_{t+1} + \gamma v_\pi(S_{t+1}) - v_\pi(S_t)]$$

with α being the stepsize. And this also applies to the q function:

$$q_\pi(S_t, A_t) \leftarrow q_\pi(S_t, A_t) + \alpha [R_{t+1} + \gamma q_\pi(S_{t+1}, A_{t+1}) - q_\pi(S_t, A_t)]$$

2.7 Policy Gradient Methods

When the state space is arbitrarily large, we cannot find an optimal policy or the optimal value function given the limited resources and time. Instead of keeping track of the value and the action(s) to select for each state (tabular methods), we can use parameterized functions, for example, artificial neural networks (ANNs), to approximate the value function and the policy function.

Previously methods learned the values of actions and then select actions based on their estimated action values (q function). Now we consider methods that instead learn a parameterized policy that select actions without consulting a value function. A value function may still be used to learn the policy parameter, but is not required for action selection.

We use θ and w to denotes the policy's parameter vector and value function's weight vector, respectively. Then we have our policy $\pi(a \mid s, \theta) = \Pr \{A_t = a \mid S_t = s, \theta_t = \theta\}$ and value function $v_\pi(s, w)$. We also define the scalar performance measure (objective) to be

$$J(\theta) = \begin{cases} v_{\pi_\theta}(s_0); & \text{episodic case} \\ \lim_{t \rightarrow \infty} \mathbb{E} [R_t \mid S_0, A_{0:t-1} \sim \pi]; & \text{continuing case} \end{cases}$$

Policy gradient methods seeks to maximize this performance measure, so their updates approximate gradient ascent in J :

$$\theta_{t+1} = \theta_t + \alpha \widehat{\nabla J(\theta_t)}$$

where $\widehat{\nabla J(\theta_t)}$ is a stochastic estimate whose expectation approximates the gradient of the performance measure w.r.t θ_t .

2.7.1 The Policy Gradient Theorem. The policy gradient theorem says

$$\nabla J(\theta) \propto \sum_s \mu(s) \sum_a q_\pi(s, a) \nabla \pi(a|s, \theta)$$

where μ is the stationary distribution of the state. Thus we can derive

$$\begin{aligned} \nabla J(\theta) &\propto \sum_s \mu(s) \sum_a q_\pi(s, a) \nabla \pi(a|s, \theta) \\ &= \mathbb{E}_\pi \left[\sum_a q_\pi(S_t, a) \nabla \pi(a|S_t, \theta) \right] \\ &= \mathbb{E}_\pi \left[\sum_a \pi(a|S_t, \theta) q_\pi(S_t, a) \frac{\nabla \pi(a|S_t, \theta)}{\pi(a|S_t, \theta)} \right] \\ &= \mathbb{E}_\pi \left[q_\pi(S_t, A_t) \frac{\nabla \pi(A_t|S_t, \theta)}{\pi(A_t|S_t, \theta)} \right] \end{aligned}$$

The policy gradient theorem can be generalized to include a comparison of the action value to an arbitrary baseline $b(s)$:

$$\nabla J(\theta) \propto \sum_s \mu(s) \sum_a (q_\pi(s, a) - b(s)) \nabla \pi(a|s, \theta)$$

with the same derivation we have

$$\nabla J(\theta) \propto \mathbb{E}_\pi \left[(q_\pi(S_t, A_t) - b(S_t)) \frac{\nabla \pi(A_t|S_t, \theta)}{\pi(A_t|S_t, \theta)} \right]$$

And the discounted version of above is:

$$\nabla J(\theta) \propto \mathbb{E}_\pi \left[\gamma^t (q_\pi(S_t, A_t) - b(S_t)) \frac{\nabla \pi(A_t|S_t, \theta)}{\pi(A_t|S_t, \theta)} \right]$$

One natural choice for baseline is an estimate of the state value, $\hat{v}(S_t, \mathbf{w})$, where \mathbf{w} is the weight vector to be learned.

2.7.2 Actor-Critic Methods. Recall the TD method, we make use of $q_\pi(s, a) = \mathbb{E}[R_{t+1} + \gamma v_\pi(S_{t+1}) | S_t = s, A_t = a]$, by taking one sample estimate, and choosing $\hat{v}(S_t, \mathbf{w})$ as baseline, we have

$$\nabla J(\theta) \propto \mathbb{E}_\pi \left[(R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}) - \hat{v}(S_t, \mathbf{w})) \frac{\nabla \pi(A_t|S_t, \theta)}{\pi(A_t|S_t, \theta)} \right]$$

Again, by taking one sample estimate we have the update formula for θ :

$$\begin{aligned} \theta_{t+1} &= \theta_t + \alpha_\theta (R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}) - \hat{v}(S_t, \mathbf{w})) \frac{\nabla \pi(A_t|S_t, \theta_t)}{\pi(A_t|S_t, \theta_t)} \\ &= \theta_t + \alpha_\theta \delta_t \nabla \ln \pi(A_t|S_t, \theta_t) \end{aligned}$$

As for the weight vector \mathbf{w} for the estimated value function, at each time step t , we want to update it by minimizing the squared error, $[v_\pi(S_t) - \hat{v}(S_t, \mathbf{w}_t)]^2$:

$$\begin{aligned} \mathbf{w}_{t+1} &= \mathbf{w}_t - \alpha'_w \nabla [v_\pi(S_t) - \hat{v}(S_t, \mathbf{w}_t)]^2 \\ &= \mathbf{w}_t - \alpha_w (v_\pi(S_t) - \hat{v}(S_t, \mathbf{w}_t)) \nabla \hat{v}(S_t, \mathbf{w}_t) \end{aligned}$$

where v_π is the true value function. Again, since $v_\pi(s) = \mathbb{E}[R_{t+1} + \gamma v_\pi(S_{t+1}) | S_t = s]$, we take one sample estimate, which gives us the update formula for \mathbf{w} :

$$\begin{aligned} \mathbf{w}_{t+1} &= \mathbf{w}_t - \alpha_w (R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}) - \hat{v}(S_t, \mathbf{w}_t)) \nabla \hat{v}(S_t, \mathbf{w}_t) \\ &= \mathbf{w}_t - \alpha_w \delta_t \nabla \hat{v}(S_t, \mathbf{w}_t) \end{aligned}$$

The actor-critic method is based on the above updates formulas, details are shown in Algorithm 1.

Algorithm 1 ACTOR-CRITIC

- 1: Input: a differentiable policy parameterization $\pi(a|s, \theta)$
 - 2: Input: a differentiable state-value function parameterization $\hat{v}(s, \mathbf{w})$
 - 3: Step sizes $\alpha_\theta > 0, \alpha_w > 0$
 - 4: $\theta \leftarrow \mathbf{0}, \mathbf{w} \leftarrow \mathbf{0}$
 - 5: **loop** (for each episode)
 - 6: Initialize S (first state of the episode)
 - 7: **loop** (for each time step i)
 - 8: $A \sim \pi(\cdot|S, \theta)$
 - 9: Take action A , observe S', R
 - 10: $\delta \leftarrow R + \gamma \hat{v}(S', \mathbf{w}) - \hat{v}(S, \mathbf{w})$
 - 11: $\mathbf{w} \leftarrow \mathbf{w} + \alpha_w \delta \nabla \hat{v}(S, \mathbf{w})$
 - 12: $\theta \leftarrow \theta + \alpha_\theta \gamma^i \delta \nabla \ln \pi(A|S, \theta)$
 - 13: $S \leftarrow S'$
 - 14: **end loop**
 - 15: **end loop**
-

3. BUILD THE ENVIRONMENT

The environment does the following things:

1. record the current state that can be queried anytime
2. when an action, i.e. a price, is given according to the current state, it outputs the reward, and goes to next state

At each time step, the state consists of a randomly generated customer's features and the current portfolio. After the agent performs an action, i.e. promotes a price to the customer, the reward is then computed based on the profit on this customer and the portfolio up to now.

3.1 Customer Features

A customer's features is represented by an array (or vector), denoted by x with 16 entries with each generated randomly from a distribution, including normal distribution, binomial distribution, and so on. For example, a feature vector x is described as Table 1

TABLE 1: CUSTOMER FEATURE VECTOR DESCRIPTION

gender	age	car cost	miles	brand	random feature 0	random feature 1	...
1.0	42.723	94032.096	30086.311	29.264	8.0	68.740	...

After a price, denoted by c , is promoted to a customer, the customer will give a response, denoted by r , which is 0 if he does not buy the product, and 1 otherwise. And the profit we have for that customer, denoted by pf , can be computed by

$$pf = r * (c - g(x))$$

where $g(x)$ is the cost on that customer.

3.2 Response

The response is generated by a generalized linear model, whose input is a vector of a customer features vector x appended with a price c , i.e. $r = glm(x[0], \dots, x[15], c)$.

The GLM used in the project is fitted using a randomly generated data set with each observation is a customer's features followed by a price, then followed by a response, as illustrated by Table 2.

And the link function for the GLM is chosen to be the logistic function.

3.3 State

The state at start of each time step is the feature vector of the new customer to promote (might be modified a bit), appended by the portfolio, which is also a data vector consists of the following entries:

- average profit over all customers so far, expect for the new customer at current time step, denoted by avg_pf (initialized to 0), which is computed by

$$avg_pf_t = \frac{1}{t} \sum_{i=0}^{t-1} pf_i, t > 0$$

- the portion of the buyers over all customers, expect for the new one, denoted by p (initialized to 0), which is computed by

$$p_t = \frac{1}{t} \sum_{i=0}^{t-1} r_i, t > 0$$

- the variance of the portion of buyers in their categories, denoted by var : we divide customers (expect for the new one at current time step) into 4 categories, according to their features, and for each category, we compute the portion of the customers who buy our product, denoted by pt_i with $i = 1, 2, 3, 4$, then we compute the variance of the 4 portions:

$$var = Var(pt) = \frac{\sum_{i=1}^4 (pt_i - \bar{pt})^2}{3}$$

Later we use this to compute reward to make sure the agent will promote to all types of customers instead of only focusing on one type. Note that $var \leq 1$.

Thus, the state s , or s_t , end up with the form $s_t = (x'_t[0], \dots, x'_t[-1], avg_pf_t, p_t, var_t)$, where x'_t is the modified x_t , including preprocess of converting some categorical data to meaningful data to the model.

3.4 Reward

The reward at each time step t , after the price c_t is promoted to the customer x_t , denoted by R_t , is calculated by

$$R_t = pf_t (1 - h(var_t))$$

where h is a function that indicates how much we care about the variance, i.e. how important it is to try to promote to customers of all categories. In this project I choose h to be such that $h(var) = \sqrt{var}/2$.

4. BUILD THE MODEL

With the environment well built, we now build the model for our policy and the value function. In this project we use artificial neuro networks (ANNs) as the parameterized policy and value function, each with three linear layers, and with the ReLU function as the activation function. And the input of both function is the state, as illustrated in Figure 2. where $n(S)$ denotes the vector length of the state and $n(A)$ denotes the number of actions.

We generate the policy distribution by scoring each action. E.g. denoted the policy function output by y , then the probability of choosing action i ($i = 0, 1, \dots, n(A) - 1$) is computed by

$$P(i) = \frac{y_i}{\sum_{j=0}^{n(A)-1} y_j}$$

Note that we add a Threshold layer at the end of policy function, which is defined by

$$Threshold(0.1, 0.1) = \max(0.1, x)$$

which makes sure that

1. the probabilities are valid (no negative values, greater score results in greater probability);
2. all actions have the probability to be chosen at early stage, which is beneficial for training since it increase the exploration for different actions.

5. RL WITH ACTOR-CRITIC

In this project, we consider the episodic learning procedure, where each episode has finite time steps, T , to be specified as a constant, and we will also reset the environment at the start of each episode.

5.1 The Algorithm

Algorithm 2 uses RL with actor-critic method, and with the help of what we have built in previous sections, to solve the problem. Figure 3 illustrates the algorithm.

TABLE 2: CUSTOMER FEATURE VECTOR DESCRIPTION

	customer feature 0	...	customer feature 15	price	resopnse
1	1.0	...	2.0	1109.504	1.0
2	0.0	...	1.0	665.8	0.0
...




FIGURE 2: ANN MODEL FOR POLICY AND VALUE FUNCTION

Algorithm 2 ACTOR-CRITIC FOR OUR FINANCIAL PROBLEM

- 1: Initialize ANN policy parameterization $\pi(c|s, \theta)$, with any θ
- 2: Initialize ANN state-value function parameterization $\hat{v}(s, w)$ with any w .
- 3: Step sizes $\alpha_\theta > 0, \alpha_w > 0$
- 4: **loop** (for each episode)
- 5: Reset the environment
- 6: Generate an x (first customer of the episode)
- 7: Initialize S (first state of the episode)
- 8: **for** $t = 0, 1, \dots, T$ **do**
- 9: $c \sim \pi(\cdot|S, \theta)$
- 10: Take action c , observe S', R from environment
- 11: $\delta \leftarrow R + \gamma \hat{v}(S', w) - \hat{v}(S, w)$
- 12: $w \leftarrow w + \alpha_w \delta \nabla \hat{v}(S, w)$
- 13: $\theta \leftarrow \theta + \alpha_\theta \gamma^t \delta \nabla \ln \pi(A|S, \theta)$
- 14: $S \leftarrow S'$
- 15: **end for**
- 16: **end loop**

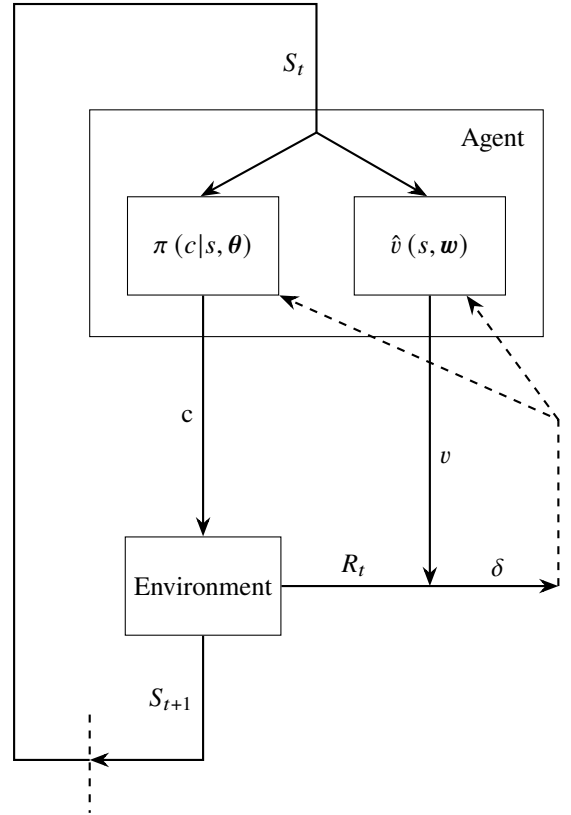


FIGURE 3: ACTOR-CRITIC FOR FINANCE ALGORITHM ILLUSTRATION

5.2 Training Result

Figure 4 shows the training result of Algorithm 1 (300 moving average reward vs iteration), with 600 iterations (5 episodes per iteration), $T = 300$, $\gamma = 0.99$, $\alpha_\theta = \alpha_w = 3 \times 10^{-4}$. The graph shows that the moving average reward converges within only 50 iterations, that is, 250 episode, 75000 parameter updates for each network, and stays quite stable afterwards.

6. ACTOR-CRITIC, REINFORCE, AND PPO

6.1 Algorithms

To further show how Actor-Critic algorithm performs, we compare it to another two popular algorithms: Reinforce (Algorithm 3) and PPO (Algorithm 4).

Algorithm 3 REINFORCE FOR OUR FINANCIAL PROBLEM

```
1: Initialize ANN policy parameterization  $\pi(c|s, \theta)$ , with any  $\theta$ 
2: Step sizes  $\alpha_\theta > 0$ 
3: loop (for each episode)
4:   Reset the environment
5:   Generate an episode  $S_0, C_0, R_1, \dots, S_{T-1}, C_{T-1}, R_T$ , following  $\pi(\cdot|\cdot, \theta)$ 
6:   for  $t = 0, 1, \dots, T$  do
7:      $G \leftarrow \sum_{k=t+1}^T \gamma^{k-t-1} R_k$ 
8:      $\theta \leftarrow \theta + \alpha_\theta \gamma^t G \nabla \ln \pi(C_t|S_t, \theta)$ 
9:   end for
10: end loop
```

6.2 Training result

To make better comparison, the policy function of the two algorithms is implemented the same way as we did in Actor-Critic algorithm, and the same is true for the value function used in PPO method.

Figure 5 shows the training result of these three algorithms (300 moving average reward vs iteration), with 600 iterations (5 episodes per iteration), $T = 300$, $\gamma = 0.99$, $\alpha_\theta = 3 \times 10^{-4}$, $\alpha_w = 3 \times 10^{-10}$ (if applicable), $\epsilon = 0.5$ (if applicable). As is shown in Figure 5, the Actor-Critic algorithm converges the fastest, within 50 iterations, followed by Reinforce algorithm, which converges within about 250 iterations, despite a small drop of moving average around iteration 300, which is quickly corrected by the algorithm within 50 iterations. PPO algorithm converges the slowest, with about 300 iterations. Despite of the speed of convergence, the moving average rewards of convergence for these three algorithms are nearly the same (around 400), this is as expected.

7. CONCLUSION

Provide a brief conclusion (3-4 lines).

ACKNOWLEDGMENTS

Place any acknowledgments here.

REFERENCES

- [1] Ning, Xiang and Lovell, Mary Rose. "On the Sliding Friction Characteristics of Unidirectional Continuous FRP Deposits." *ASME Journal of Tribology* Vol. 48 No. 5 (2002): pp. 2000–2008. DOI [10.1115/1.4042912](https://doi.org/10.1115/1.4042912).
- [2] Gibson, Tom A. and Tucker, Matthew A. *The Big Book of Cellular Studies*. John Wiley and Sons, New York (2008).
- [3] Stevens, Thomas T. "Stochastic Fields and Their Digital Simulation." *Stochastic Methods*. Martinus Publishers, Dordrecht (1999): pp. 22–36.
- [4] Wions, Thomas and Mills, Christopher D. "Structural Dynamics in Parallel Manipulation." *Proceedings of the ASME IDETC/CIE*. DETC2005-99532: pp. 777–798. New Orleans, LA, September 10–13, 2005. DOI [10.1115/1.4042912](https://doi.org/10.1115/1.4042912). URL <https://doi.org/10.1115/1.4042912>.
- [5] Smith, Robert. "Conformal Lubricated Contact of Cylindrical Surfaces Involved in a Non-Steady Motion." Ph.D. Thesis, Rochester Institute of Technology, Rochester, NY. 2002. DOI [10.1115/1.4042912](https://doi.org/10.1115/1.4042912). URL <https://doi.org/10.1115/1.4042912>.
- [6] Smith, Robert. "Conformal Lubricated Contact of Cylindrical Surfaces Involved in a Steady Motion." Master's Thesis, Rochester Institute of Technology, Rochester, NY. 1998.
- [7] Smith, Robert. "Cylindrical Surfaces Involved in a Steady Motion." BS Thesis, Rochester Institute of Technology, Rochester, NY. 1996.
- [8] Watson, David. "Evaporative Heat Transfer in the Contact Line of a Mixture." Technical Report No. HTL-26, CFD-4. Iowa State University, Ames, IA. 1982. DOI [10.1115/1.4042912](https://doi.org/10.1115/1.4042912). URL <https://doi.org/10.1115/1.4042912>.
- [9] Reynolds, William C., Kays, William M. and Kline, Stephen J. "Heat Transfer in the Incompressible Turbulent Boundary Layer. I—Constant Wall Temperature." Memorandum 12-1-58W. NASA, Washington, DC. 1958. URL <https://hdl.handle.net/2060/19980228020>.
- [10] Edwards, Donald K. "Thermal Radiation Measurements." Eckert, Ernst R. G. and Goldstein, Richard J. (eds.). *Measurement Techniques in Heat Transfer*. Vol. 130 of *AGARDograph*. Technivision Services, Slough, England (1970): Chap. 9, pp. 353–396. Published by the Advisory Group for Aerospace Research and Development of NATO.
- [11] Lienhard, John H., V. "Exterior Shape Factors from Interior Shape Factors." *ASME J. Heat Transfer* Vol. 141 No. 6 (2019): 061301. DOI [10.1115/1.4042912](https://doi.org/10.1115/1.4042912).
- [12] Braams, Johannes L. and Bezos, Javier. "babel: Localization and internationalization." Version 3.52. Comprehensive T_EX Archive Network (2020). Accessed January 3, 2021, URL <https://ctan.org/pkg/babel>.
- [13] Lienhard, John H., V. "Language Support in asmeconf: non-latin alphabets, Lua_T_EX, and fontspec." Version 1.0. Comprehensive T_EX Archive Network (2021). Accessed January 18, 2021, URL <https://www.ctan.org/pkg/ameconf>.




FIGURE 4: TRAINING RESULT FOR ALGORITHM A2C

Algorithm 4 PPO FOR OUR FINANCIAL PROBLEM

- 1: Initialize ANN policy parameterization $\pi(c|s, \theta)$, with any θ
 - 2: Initialize ANN state-value function parameterization $\hat{v}(s, w)$ with any w .
 - 3: Step sizes $\alpha_\theta > 0, \alpha_w > 0$
 - 4: **for** iteration $k = 0, 1, 2, \dots$ **do**
 - 5: Collect set of trajectories $\mathcal{D}_k = \{\tau_i\}$, following $\pi(\cdot|\cdot, \theta_k)$, where $\tau = S_0, C_0, R_1, \dots, S_{T-1}, C_{T-1}, R_T$
 - 6: **for** $t = 0, 1, \dots, T$ in each trajectory τ in \mathcal{D}_k **do**
 - 7: $G_t \leftarrow \sum_{k=t+1}^T \gamma^{k-t-1} R_k$
 - 8: $A_t \leftarrow G_t - \hat{v}(S_t)$
 - 9: **end for**
 - 10: **loop** (update multiple times)
 - 11: $\theta \leftarrow \arg \max_{\theta} \frac{1}{|\mathcal{D}_k|T} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \min \left(\frac{\pi(c_t|S_t, \theta)}{\pi(c_t|S_t, \theta_k)} A_t, \text{clip} \left(\frac{\pi(c_t|S_t, \theta)}{\pi(c_t|S_t, \theta_k)}, 1 - \epsilon, 1 + \epsilon \right) A_t \right)$
 - 12: $w \leftarrow \arg \min_w \frac{1}{|\mathcal{D}_k|T} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \left(\hat{v}(S_t, w) - G_t \right)^2$
 - 13: **end loop**
 - 14: **end for**
-




FIGURE 5: TRAINING RESULT COMPARISON FOR THREE ALGORITHM

- [14] Fourier, Jean Baptiste Joseph. *Théorie Analytique de la Chaleur*. Firmin Didot, Père et Fils, Paris (1822). URL <https://books.google.com/books?id=1TUVA AAAQAAJ>.
- [15] Johnson, Howard (ed.). *Title of Edited Book*. John Wiley and Sons, Inc., New York, NY (1980).
- [16] Yankovich, Al and Stern, Howard (eds.). *The Title of the Proceedings*, Vol. 15. Boston, MA, April 27–28. Organization (2015). URL <http://web.mit.edu>. Optional note.
- [17] Govindan, Prakash Narayan, Lam, Steven and St. John, Maximus G. “Systems Including a Condensing Apparatus Such As a Bubble Column Condenser.” US Patent #9700811 (2017). URL <https://patents.google.com/patent/US9700811B2>.
- [18] Fong, Chamberlain. “Analytical Methods for Squaring the Disc.” *ArXiv e-prints* (2015). URL <https://arxiv.org/abs/1509.06344>.
- [19] Galassi, Mark, Davies, Jim, Theiler, James, Gough, Brian, Jungman, Gerard, Alken, Patrick, Booth, Michael, Rossi, Fabrice and Ulerich, Rhys. *GNU Scientific Library*, Release 2.4. Free Software Foundation, Boston, MA (2017). URL <https://www.gnu.org/software/gsl/>.
- [20] American Mathematical Society and LATEX3 Project. “User’s Guide for the amsmath Package.” Version 2.1. Comprehensive T_EX Archive Network (2018). Accessed July 5, 2019, URL <https://ctan.org/tex-archive/macros/latex/required/amsmath/amslatex/doc.pdf>.
- [21] Lienhard, John H., V. “Wide Equations in asmeconf.cls.” Version 1.0. Comprehensive T_EX Archive Network (2022). Accessed January 10, 2022, URL <https://www.ctan.org/pkg/asmconf>.
- [22] Sharpe, Michael. “New TX Font Package.” Version 1.640. Comprehensive T_EX Archive Network (2020). Accessed January 4, 2021, URL <https://ctan.org/pkg/newtx>.
- [23] Sharpe, Michael. “The mathalfa Package.” Version 1.10. Comprehensive T_EX Archive Network (2018). Accessed April 27, 2019, URL <https://ctan.org/pkg/mathalfa>.
- [24] Tolušis, Sigitas. “The flushend Package.” Version 3.3. Comprehensive T_EX Archive Network (2017). Accessed October 1, 2019, URL <https://www.ctan.org/pkg/flushend>.
- [25] Böttcher, Stephan I. and Lück, Uwe. “lineno.sty: A L^AT_EX Package to Attach Line Numbers to Paragraphs.” Version 4.41. Comprehensive T_EX Archive Network (2005). Accessed August 17, 2019, URL <https://www.ctan.org/pkg/lineno>.
- [26] van Oostrum, Piet. “The fancyhdr and extramarks Packages.” Version 3.10. Comprehensive T_EX Archive Network (2019). Accessed August 17, 2019, URL <https://www.ctan.org/pkg/fancyhdr>.
- [27] Fairbairns, Robin, Rahtz, Sebastian and Barroca, Leonor. “A Package for Rotated Objects in L^AT_EX.” Version 2.16d. Comprehensive T_EX Archive Network (2016). Accessed October 2, 2019, URL <https://www.ctan.org/pkg/rotating>.
- [28] van der Walt, Stéfan and Smith, Nathaniel. “mpl colormaps.” GitHub, San Francisco, CA (2015). Accessed August 26, 2018, URL <http://bids.github.io/colormap/>.
- [29] Churchill, Winston S. *The Gathering Storm*, 1st ed. Houghton Mifflin Co., Boston (1948): Chap. 5, pp. 66–89. “The Locust Years”.
- [30] Kirk, James Tiberius. *Decline and Fall of the Romulan Empire*, 7th ed. Vol. 23 of *Humankind’s Greatest Writings*. Vulcan Free Press, T’Paal (2288). URL <http://web.mit.edu>. To appear.

APPENDIX A. THE VECTOR PRODUCT $\vec{A} \times \vec{B}$

This brief illustration of an appendix shows the numbering of the appendix and equations. Equations are numbered consecutively, following those in the paper. Consider $\rho \neq \text{fn}(p)$:

$$\frac{d\Gamma}{dt} = \frac{d}{dt} \int_{\mathcal{G}} \mathbf{u} \cdot d\mathbf{r} \quad (1)$$

$$= \int_{\mathcal{G}} \frac{D\mathbf{u}}{Dt} \cdot d\mathbf{r} + \underbrace{\int_{\mathcal{G}} \mathbf{u} \cdot d\left(\frac{d\mathbf{r}}{dt}\right)}_{=0} \quad (2)$$

$$= \iint_{\mathcal{S}} \nabla \times \frac{D\mathbf{u}}{Dt} \cdot d\mathbf{A} \quad (3)$$

$$= \iint_{\mathcal{S}} \nabla p \times \nabla \left(\frac{1}{\rho}\right) \cdot d\mathbf{A} \quad (4)$$